

# Logistic Regression Cheatsheet

BAN-501: Business Analytics

## Contents

<b>1 Odds and Log-Odds</b>	<b>2</b>
1.1 From Probability to Odds . . . . .	2
1.1.1 Worked Examples . . . . .	2
1.2 From Odds to Log-Odds (The Logit Function) . . . . .	2
1.3 Probability–Odds–Log-Odds Conversion Table . . . . .	2
<b>2 The Logistic (Sigmoid) Function</b>	<b>3</b>
2.1 Key Properties . . . . .	3
2.2 Key Values . . . . .	3
2.3 Relationship Between Logit and Sigmoid . . . . .	3
<b>3 The Logistic Regression Model</b>	<b>3</b>
3.1 Model Equation . . . . .	3
3.2 Coefficient Interpretation . . . . .	4
<b>4 Worked Prediction Example</b>	<b>4</b>
<b>5 Python Code Snippets</b>	<b>4</b>
5.1 Fitting Logistic Regression with Statsmodels . . . . .	4
5.2 Making Predictions . . . . .	5
5.3 Manual Probability Calculation . . . . .	5
<b>6 Classification Metrics</b>	<b>5</b>
6.1 Confusion Matrix . . . . .	5
6.2 Key Metrics . . . . .	6
6.3 Worked Example . . . . .	6
6.4 When to Prioritize Each Metric . . . . .	7
6.5 Python Code for Metrics . . . . .	7
<b>7 ROC Curve and AUC</b>	<b>7</b>
7.1 Definitions . . . . .	7
7.2 ROC Curve . . . . .	8
7.3 AUC (Area Under the ROC Curve) . . . . .	8
7.4 Python Code for ROC Curve . . . . .	8
<b>8 Quick Reference</b>	<b>9</b>
8.1 Key Formulas . . . . .	9

8.2 Coefficient Interpretation Checklist . . . . .	10
--	----

# 1 Odds and Log-Odds

## 1.1 From Probability to Odds

Given a probability  $p = P(Y = 1)$ , the **odds** are defined as:

$$\text{odds} = \frac{p}{1-p} = \frac{P(Y = 1)}{P(Y = 0)}$$

**Interpretation:** Odds represent the ratio of success to failure. An odds of 4 means “4 to 1 in favor.”

### 1.1.1 Worked Examples

- $p = 0.2$ : odds =  $\frac{0.2}{1-0.2} = \frac{0.2}{0.8} = 0.25$  (4 to 1 against)
- $p = 0.5$ : odds =  $\frac{0.5}{1-0.5} = \frac{0.5}{0.5} = 1.0$  (even odds)
- $p = 0.8$ : odds =  $\frac{0.8}{1-0.8} = \frac{0.8}{0.2} = 4.0$  (4 to 1 in favor)

## 1.2 From Odds to Log-Odds (The Logit Function)

The **log-odds** (or **logit**) is the natural logarithm of the odds:

$$\text{logit}(p) = \log\left(\frac{p}{1-p}\right)$$

**Key insight:** The logit function maps probabilities from  $[0, 1]$  to  $(-\infty, +\infty)$ .

## 1.3 Probability–Odds–Log-Odds Conversion Table

Probability ( $p$ )	Odds	Log-Odds
0.01	0.0101	-4.60
0.10	0.1111	-2.20
0.20	0.2500	-1.39
0.25	0.3333	-1.10
0.50	1.0000	0.00
0.75	3.0000	+1.10
0.80	4.0000	+1.39
0.90	9.0000	+2.20
0.99	99.0000	+4.60

**Properties:**

- Log-odds are **symmetric** around 0
- $p = 0.5$  corresponds to log-odds = 0
- Log-odds range from  $-\infty$  to  $+\infty$

## 2 The Logistic (Sigmoid) Function

The **logistic function** (also called the **sigmoid function**) is the inverse of the logit:

$$\sigma(z) = \frac{1}{1 + e^{-z}} = \frac{e^z}{1 + e^z}$$

where  $z$  represents the log-odds (or any linear predictor).

### 2.1 Key Properties

1. **Bounded output:** Always produces values in  $(0, 1)$ —valid probabilities
2. **S-shaped:** Captures natural threshold behavior in classification
3. **Symmetric:**  $\sigma(-z) = 1 - \sigma(z)$
4. **Smooth:** Differentiable everywhere (important for optimization)

### 2.2 Key Values

Input ( $z$ )	Output $\sigma(z)$
$-\infty$	$\rightarrow 0$
-2	0.119
-1	0.269
0	0.500
+1	0.731
+2	0.881
$+\infty$	$\rightarrow 1$

### 2.3 Relationship Between Logit and Sigmoid

The logit and sigmoid are **inverse functions**:

$$\sigma(\text{logit}(p)) = p \quad \text{and} \quad \text{logit}(\sigma(z)) = z$$

## 3 The Logistic Regression Model

### 3.1 Model Equation

**Log-odds form:**

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_k x_k$$

**Probability form:**

$$P(Y = 1 | \mathbf{x}) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \cdots + \beta_k x_k)}}$$

For the simple case with one predictor:

$$P(Y = 1 | x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$

### 3.2 Coefficient Interpretation

- $\beta_0$  (intercept): The log-odds of  $Y = 1$  when all predictors equal zero
- $\beta_j$  (coefficient): The change in log-odds for a one-unit increase in  $x_j$ , holding other variables constant
- $e^{\beta_j}$  (odds ratio): The multiplicative change in odds for a one-unit increase in  $x_j$

**Interpreting the odds ratio:**

- $e^{\beta_j} > 1$ : Odds increase (positive association)
- $e^{\beta_j} = 1$ : No effect ( $\beta_j = 0$ )
- $e^{\beta_j} < 1$ : Odds decrease (negative association)

## 4 Worked Prediction Example

**Given:** A fitted logistic regression model with:

- $\beta_0 = -2$  (intercept)
- $\beta_1 = 1.5$  (coefficient for  $x$ )
- New observation:  $x = 2$

**Step 1: Compute the linear predictor (log-odds)**

$$z = \beta_0 + \beta_1 x = -2 + 1.5(2) = -2 + 3 = 1$$

**Step 2: Apply the sigmoid function to get probability**

$$P(Y = 1 \mid x = 2) = \sigma(1) = \frac{1}{1 + e^{-1}} = \frac{1}{1 + 0.368} = \frac{1}{1.368} \approx 0.731$$

**Step 3: Interpret the result**

- The predicted probability of  $Y = 1$  is approximately 73.1%
- If using a threshold of 0.5, we would classify this observation as  $\hat{Y} = 1$

**Bonus: Interpret the coefficient**

- $\beta_1 = 1.5$  means a one-unit increase in  $x$  increases log-odds by 1.5
- Odds ratio:  $e^{1.5} \approx 4.48$ , so a one-unit increase in  $x$  multiplies the odds by about 4.5

## 5 Python Code Snippets

### 5.1 Fitting Logistic Regression with Statsmodels

```
1 import numpy as np
2 import statsmodels.api as sm
3
4 # Prepare data (add constant for intercept)
```

```

5 X = sm.add_constant(X_train)
6
7 # Fit model
8 model = sm.Logit(
9     endog=y_train,
10    exog=X,
11 )
12 result = model.fit()
13
14 # View summary
15 print(result.summary())
16
17 # Extract coefficients
18 beta_0 = result.params[0] # intercept
19 beta_1 = result.params[1] # first coefficient
20
21 # Odds ratios
22 odds_ratios = np.exp(result.params)

```

## 5.2 Making Predictions

```

1 # Prepare new data
2 X_new = sm.add_constant(X_test)
3
4 # Predict probabilities
5 probabilities = result.predict(X_new)
6
7 # Convert to class predictions (threshold = 0.5)
8 predictions = (probabilities >= 0.5).astype(int)

```

## 5.3 Manual Probability Calculation

```

1 def sigmoid(z):
2     """Compute the sigmoid/logistic function."""
3     return 1 / (1 + np.exp(-z))
4
5 # Given coefficients and new x value
6 beta_0 = -2
7 beta_1 = 1.5
8 x_new = 2
9
10 # Compute probability
11 z = beta_0 + beta_1 * x_new # linear predictor
12 prob = sigmoid(z) # probability
13 print(f"P(Y=1 | x={x_new}) = {prob:.3f}")

```

# 6 Classification Metrics

## 6.1 Confusion Matrix

A confusion matrix summarizes classification results:

	Predicted Negative	Predicted Positive
Actual Negative	True Negative (TN)	False Positive (FP)
Actual Positive	False Negative (FN)	True Positive (TP)

## 6.2 Key Metrics

**Recall (Sensitivity, True Positive Rate):**

$$\text{Recall} = \frac{TP}{TP + FN}$$

“Of all actual positives, what fraction did we correctly identify?”

**Precision (Positive Predictive Value):**

$$\text{Precision} = \frac{TP}{TP + FP}$$

“Of all predicted positives, what fraction were actually positive?”

**Specificity (True Negative Rate):**

$$\text{Specificity} = \frac{TN}{TN + FP}$$

“Of all actual negatives, what fraction did we correctly identify?”

**Accuracy:**

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

**F1 Score:** Harmonic mean of precision and recall:

$$F_1 = 2 \cdot \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} = \frac{2 \cdot TP}{2 \cdot TP + FP + FN}$$

## 6.3 Worked Example

Given confusion matrix:

	Pred. Neg.	Pred. Pos.
Actual Neg.	TN = 50	FP = 10
Actual Pos.	FN = 5	TP = 35

**Calculations:**

$$\text{Recall} = \frac{35}{35 + 5} = \frac{35}{40} = 0.875$$

$$\text{Precision} = \frac{35}{35 + 10} = \frac{35}{45} \approx 0.778$$

$$\text{Specificity} = \frac{50}{50 + 10} = \frac{50}{60} \approx 0.833$$

$$\text{Accuracy} = \frac{35 + 50}{35 + 50 + 10 + 5} = \frac{85}{100} = 0.85$$

$$F_1 = 2 \cdot \frac{0.778 \times 0.875}{0.778 + 0.875} = 2 \cdot \frac{0.681}{1.653} \approx 0.824$$

## 6.4 When to Prioritize Each Metric

Metric	Prioritize when...	Example
Recall	False negatives are costly	Disease screening, fraud detection
Precision	False positives are costly	Spam filtering, criminal conviction
F1 Score	Need balance between precision and recall	General classification tasks
Accuracy	Classes are balanced and errors are equally costly	Balanced datasets

## 6.5 Python Code for Metrics

```

1 from sklearn.metrics import (
2     confusion_matrix,
3     classification_report,
4     precision_score,
5     recall_score,
6     f1_score,
7     accuracy_score,
8 )
9
10 # Compute confusion matrix
11 cm = confusion_matrix(
12     y_true=y_test,
13     y_pred=y_pred,
14 )
15 print("Confusion Matrix:")
16 print(cm)
17
18 # Individual metrics
19 print(f"Accuracy: {accuracy_score(y_test, y_pred):.3f}")
20 print(f"Precision: {precision_score(y_test, y_pred):.3f}")
21 print(f"Recall: {recall_score(y_test, y_pred):.3f}")
22 print(f"F1 Score: {f1_score(y_test, y_pred):.3f}")
23
24 # Full classification report
25 print(classification_report(y_test, y_pred))

```

## 7 ROC Curve and AUC

### 7.1 Definitions

**True Positive Rate (TPR)** = Recall = Sensitivity:

$$\text{TPR} = \frac{TP}{TP + FN}$$

**False Positive Rate (FPR)** = 1 – Specificity:

$$\text{FPR} = \frac{FP}{FP + TN}$$

## 7.2 ROC Curve

The **Receiver Operating Characteristic (ROC) curve** plots TPR vs. FPR at various classification thresholds.

### How it's generated:

1. Start with predicted probabilities for all observations
2. Vary the classification threshold from 0 to 1
3. At each threshold, compute TPR and FPR
4. Plot (FPR, TPR) pairs as a curve

### Key points on the ROC curve:

- (0, 0): Threshold = 1 (predict all negative)
- (1, 1): Threshold = 0 (predict all positive)
- (0, 1): Perfect classifier
- Diagonal line: Random classifier (no skill)

## 7.3 AUC (Area Under the ROC Curve)

**AUC** measures the overall ability of the model to discriminate between classes.

### Interpretation:

- AUC = probability that a randomly chosen positive instance is ranked higher than a randomly chosen negative instance
- AUC = 1.0: Perfect discrimination
- AUC = 0.5: No discrimination (random guessing)
- AUC < 0.5: Worse than random (predictions are inverted)

### AUC guidelines:

AUC Range	Interpretation
0.90–1.00	Excellent
0.80–0.90	Good
0.70–0.80	Fair
0.60–0.70	Poor
0.50–0.60	Fail

## 7.4 Python Code for ROC Curve

```
1 import matplotlib.pyplot as plt
2 from sklearn.metrics import roc_curve, roc_auc_score
3
4 # Get predicted probabilities (not class predictions)
5 y_proba = model.predict_proba(X_test)[:, 1]
6
```

```

7 # Compute ROC curve
8 fpr, tpr, thresholds = roc_curve(
9     y_true=y_test,
10    y_score=y_proba,
11 )
12
13 # Compute AUC
14 auc = roc_auc_score(
15     y_true=y_test,
16     y_score=y_proba,
17 )
18
19 # Plot ROC curve
20 fig, ax = plt.subplots(figsize=(6, 6))
21 ax.plot(fpr, tpr, label=f"ROC Curve (AUC = {auc:.3f})")
22 ax.plot([0, 1], [0, 1], "k--", label="Random Classifier")
23 ax.set_xlabel("False Positive Rate")
24 ax.set_ylabel("True Positive Rate")
25 ax.set_title("ROC Curve")
26 ax.legend()
27 plt.show()

```

## 8 Quick Reference

### 8.1 Key Formulas

Concept	Formula
Odds	$\frac{p}{1-p}$
Logit (log-odds)	$\log\left(\frac{p}{1-p}\right)$
Sigmoid	$\sigma(z) = \frac{1}{1+e^{-z}}$
Logistic regression	$P(Y=1) = \frac{1}{1+e^{-(\beta_0+\beta_1x)}}$
Odds ratio	$e^{\beta_j}$
Recall	$\frac{TP}{TP+FN}$
Precision	$\frac{TP}{TP+FP}$
F1 Score	$\frac{2 \cdot TP}{2 \cdot TP + FP + FN}$
TPR (for ROC)	$\frac{TP}{TP+FN}$
FPR (for ROC)	$\frac{FP}{FP+TN}$

## 8.2 Coefficient Interpretation Checklist

1. Identify  $\beta_j$  from the model output
2. **Log-odds interpretation:** “A one-unit increase in  $x_j$  changes the log-odds by  $\beta_j$ ”
3. Calculate odds ratio:  $e^{\beta_j}$
4. **Odds ratio interpretation:** “A one-unit increase in  $x_j$  multiplies the odds by  $e^{\beta_j}$ ”
5. If  $e^{\beta_j} > 1$ : odds increase; if  $e^{\beta_j} < 1$ : odds decrease